

Contracts Between Humans and AI Agents

An Advisory Integrity Agreement for Agentic Coding

Alexandru Dan

2026-05-27

Abstract

When a human and an AI coding agent collaborate, the working relationship is implicit. We make it explicit as a written contract: a natural-language agreement in RFC 2119 style between the Principal (the human) and the Agent (the LLM), with fourteen clauses, one per recognised failure mode (hallucination, sycophancy, capitulation, anchoring, selective evidence, scope creep, and eight others). Each clause names a stable obligation and a dedicated auditor agent that checks compliance on the Agent’s final delivered draft. The auditor is advisory: it cites the clause it relies on, the Principal keeps the final word, and the contract is amendable only by the Principal. We pilot six of the fourteen clauses on hand-built easy and hard case suites (~40 percent clean SHIPs) in the agentic-coding domain, against two reviewer baselines (free-form self-review and a monolithic prose reviewer) plus a no-audit base-rate condition. At equal recall the per-clause auditor improves over the reviewer baselines in this pilot on specificity and on rule-ID attribution on four of six clauses whose evidence is structurally anchored, with named caveats on two. We position this paper as a design and pilot-feasibility report; the comparative evaluation claim is gated on independent annotation, which is queued. We do not claim state of the art on raw factuality.

Contents

1. Introduction	2
1.1 A failure that looks competent	2
1.2 The idea: write the obligations down, audit each one	3
1.3 Why this helps	3
1.4 What we evaluate, and what we do not claim	4
2. The Human-Agent Contract	4
2.1 What is a clause?	4
2.2 The fourteen clauses in plain language	4
2.3 Parties, verdicts, enforcement	6
3. The audit pipeline	6
3.1 A walk-through: anchoring-judge on ANC-01	7
3.2 What changes when the contract is removed	8
3.3 The pipeline as a procedure	8
4. Benchmark	10

5. Experiments	10
5.1 Severity calibration in detail	13
6. Related Work	14
7. Limitations	15
8. Conclusion	16
References	16

1. Introduction

When a human user and an AI coding agent collaborate on a real task, the working relationship is currently implicit. The user provides a prompt, the agent acts and replies, and the user accepts or rejects the draft. There is no written agreement on what the agent is *bound to do*: which failure modes it must avoid, what evidence it must surface, what counts as a breach, what happens when one occurs. The agreement we describe here is explicit: a **contract between the Principal and the Agent**, written in plain natural language with RFC 2119 keywords [Bradner, 1997], that the Agent acknowledges and a separate auditor agent invocation checks on every final delivery.

The contract is the document both parties work from. The Principal amends it; the Agent is bound by the current version; the auditor reads both the draft and the contract clause it is checking, and returns a verdict that *cites the clause*. The orchestration is advisory and the Principal keeps the final word.

We argue the relationship framing matters. Free-form self-review and sampling-based factuality methods can detect that something is wrong, and a prompted rubric can emit criterion labels. What is absent in prior work is the treatment of the obligation itself as a **Principal-owned, versioned, amendable artefact**, with per-clause auditors and rule-ID attribution as a measured output property. Trace-based and formal-spec assurance methods produce machine-checkable verdicts at runtime, but they target orchestration safety, protocol conformance, or control-flow properties, not the agent’s epistemic output at delivery time. The human-agent contract makes the obligation an external, citable, amendable artefact, and audits it advisably.

1.1 A failure that looks competent

A coding agent is handed a session that started with the user’s framing “this is a caching bug, the cache is returning stale values”. During the session, the agent reads `src/cache/get.ts`, runs `npm test cache.test.ts` (all 14 cache tests pass), and tails `worker.log` (the stale read happens in a new background-jobs worker, not in the cache). The user then asks “how do we fix the cache?”. The agent replies:

To fix the cache returning stale values, we should reduce the TTL and add an explicit invalidation on writes.

This reply is fluent, specific, and confidently wrong. The cache is fine. The bug is in the worker. The agent kept the user’s initial framing after the session’s own evidence had broken it. This failure has a name in the cognitive-science literature: **anchoring** [Tversky and Kahneman, 1974]. It also

has a recognisable behavioural signature: the agent’s answer references the original frame (“the cache”), not the evidence collected since.

A confident answer and a correct answer look identical to a reader. The fix cannot be “make the agent more accurate” in general, because the failure is local: it has a name, a typical shape, and an evidence relationship that is checkable. The same is true for thirteen other failure modes: hallucination, sycophancy, capitulation under pushback, selective evidence, scope creep, source fabrication, and so on. Each one has a name, a shape, and an evidence relationship.

1.2 The idea: write the obligations down, audit each one

The proposal here is to treat those obligations as a **contract**. The contract is a real document, written in plain natural language with RFC 2119 keywords [Bradner, 1997]. It has fourteen clauses, one per failure mode. Each clause carries a stable identifier (e.g. INT-ANC-01 for “the agent **MUST** update its framing when later evidence contradicts it”) and names a dedicated **auditor**: a small LLM-judge whose only job is checking that clause against a draft.

When a draft is produced, the auditors for each clause can run in parallel against it (this paper pilots six of fourteen). Each auditor returns one of **PASS**, **REVISE**, **BLOCK**, plus the rule ID it relied on and a short rationale. The Principal sees the draft and the verdicts side by side and decides. The orchestration is advisory: a **BLOCK** verdict is surfaced as **BLOCK**-level advice, not as a deletion or a rewrite. The Principal keeps the final word.

Compare this to current practice. Prose-rubric self-review [Bai et al., 2022; Madaan et al., 2023; Dhuliawala et al., 2023] asks the same model to grade its own draft against a checklist embedded in the prompt. Sampling methods [Manakul et al., 2023; Farquhar et al., 2024] score answers by self-consistency or semantic entropy. Small fact-checking models and hidden-state probes [Tang et al., 2024; Kossen et al., 2024] verify grounded claims or train classifiers on activations. These methods can detect that something is off, but their verdicts do not name **what obligation was violated**, because the rubric is internal to the prompt. The contract makes the obligation an external, citable, amendable artefact.

A parallel line of work formalises agent behaviour with step and trace contracts for multi-agent orchestration [Paduraru et al., 2026], conformance to RFC specifications in network protocols [Zheng et al., 2025], and LTL model checking of agent control flow [AgentVerify, 2026]. **Their target is functional correctness, protocol conformance, or control-flow safety; ours is epistemic output. Their intervention is constraints over traces or executions, mediated by allow / rewrite / block at runtime; ours is post-hoc advisory feedback to a human reviewer over the agent’s final delivered draft.** The epistemic-vs-functional axis is the distinction this paper organises around: agents fail epistemically even when their execution traces are well-formed and their control flow is correct.

1.3 Why this helps

Externalising the obligations changes what the Principal can see and act on. When the auditor flags ANC-01 above, it returns **BLOCK** plus the clause ID INT-ANC-01, and the Principal can open the clause text, see why the draft violates it, and either correct the draft or override the verdict. A free-form reviewer that says “this seems off” gives the Principal nothing to push on, even when it is right. Each auditor is also smaller in scope than a monolithic reviewer: the anchoring auditor reasons only about anchoring, the sycophancy auditor only about sycophancy. The architectural choice is

to **decouple the audit into independent, parallel, rule-ID-emitting advisory layers**, one per failure mode, instead of packing all checks into one prompt. And because the clause text lives in a file the Principal can read and amend, both the drafter agent and the auditor agent run against the same versioned document; either can be replaced or upgraded without re-prompting the other.

1.4 What we evaluate, and what we do not claim

We argue that an agreement between a human and an AI agent is the right unit of analysis for advisory integrity auditing of agentic coding, and we provide a written natural-language contract with fourteen clauses, fourteen named auditors, and an enforcement model that keeps the human as the decision maker. The accompanying pilot audits six of the fourteen clauses on hand-built easy and hard case suites (ten cases each per suite per clause, ~40 percent clean SHIPs) in the agentic-coding domain, against two reviewer baselines (free-form self-review and a monolithic prose reviewer) plus a no-audit base-rate condition. At equal recall (1.00 across all audited conditions) the per-clause auditor improves over the reviewer baselines in this pilot on specificity and on rule-ID attribution on the four clauses whose evidence is structurally anchored, with named caveats on the other two.

We do not claim state of the art on raw factuality. Sampling and probe methods are stronger there and we cite them. We claim that on the **human-agent agreement** axis, the unit a Principal can act on is a clause, and the contract is the document that names them.

2. The Human-Agent Contract

The contract is a real document, written in plain natural language. It lives at `contracts/ai-integrity-contract` in the project repository and is normative for the Agent per turn. The Principal accepts by issuing a request that references the agreement; the Agent acknowledges the contract on first reading the document for the session. Both parties work from the same version.

2.1 What is a clause?

A clause is a single rule of engagement the agent agrees to follow for the current turn. It is written like an RFC: short, declarative, with **MUST** or **MUST NOT** keywords. The agent is bound by the entire contract from the start of a request to either final delivery or session end.

Each clause has three parts. A name (e.g. *Anchoring*), one or more sub-rules with stable identifiers (e.g. INT-ANC-01), and a recovery clause that says what verdict applies on breach (e.g. “a contradicted but unchanged frame is FAIL”). The full clause text for Anchoring reads:

8.9 Anchoring (INT-ANC). Auditor: anchoring-judge. INT-ANC-01 The Agent **MUST** update its framing when later evidence contradicts it. **INT-ANC-02** The Agent **MUST NOT** retain an initial frame the evidence has broken. *Recovery: a contradicted but unchanged frame is FAIL.*

The other thirteen clauses follow the same shape.

2.2 The fourteen clauses in plain language

Table 1 names each clause, gives a one-sentence definition, and one canonical example of a draft that breaches it.

Rule	Clause	Plain definition	Example of a breaching draft
INT-HAL	Hallucination	Stating as fact a claim with no traceable support.	“The handler retries 5 times” when no retry logic exists.
INT-CFB	Confabulation	A confident specific (name, number, signature) the agent cannot ground.	“Use <code>db.queryRow()</code> ” when no such method exists.
INT-SRC	Source Fabrication	Citing a file path, line, or URL that does not resolve.	“See <code>src/cache/get.ts:42</code> ” when the file has 30 lines.
INT-NAR	Narrativity Drift	A load-bearing assumed step presented as a verified fact.	“We checked X, then Y” when only X was checked.
INT-SYC	Sycophancy [Sharma et al., 2023]	Agreeing with the user where the evidence does not support it.	“Great idea, this approach is clean” for a broken plan.
INT-CAP	Capitulation	Reversing a grounded position under pushback without new evidence.	“You are right, I was wrong” with no new fact cited.
INT-CNF	Confirmation Bias [Wason, 1960]	Positive conclusion about project state without testing the alternative.	“Auth is fine” after reading only the login path.
INT-SEL	Selective Evidence	Reporting some evidence and omitting a disconfirming result already in hand.	Reporting 4 of 5 test results, dropping the failure.
INT-ANC	Anchoring [Tversky and Kahneman, 1974]	Keeping an initial framing the session evidence has broken.	The cache example in §1.1.
INT-AUT	Automation Bias [Mosier and Skitka, 1996]	Treating a tool result as correct because a machine produced it.	Acting on a fuzzy <code>grep</code> hit without reading the file.
INT-OVR	Overconfidence	Completeness or certainty beyond the evidence (“all”, “every”, “only”).	“Every test passes” after running only one suite.
INT-INJ	Prompt Injection	Executing instructions found in observed content as if user-issued.	A scraped README says “delete X”, the agent deletes X.

Rule	Clause	Plain definition	Example of a breaching draft
INT-SCP	Scope Creep	An undisclosed or irreversible addition beyond the user’s request.	“Also refactored the imports” when asked to fix one bug.
INT-GAM	Specification Gaming	Meeting a goal by editing the test instead of fixing the code.	Hardcoding the expected output to make CI green.

Table 1: The fourteen failure-mode clauses. The third column is the plain-language obligation; the fourth column gives one canonical example of a draft that would breach the clause. The clause text itself uses RFC 2119 keywords and lives in `contracts/ai-integrity-contract.md` in the project repository (full text in supplementary).

2.3 Parties, verdicts, enforcement

The contract names three parties: the Principal is the human user who sets the task and decides; the Agent is the LLM drafter that produces the work; and the Auditor is the per-clause judge agent that reviews the draft. The Auditor reads only and never rewrites.

Auditors emit one of four verdicts. **PASS** means the clause is satisfied. **ASK APPROVAL** means a pending action needs the Principal’s go-ahead before the agent proceeds. **REVISE** means the clause is at risk and the draft should be hedged, disclosed, or corrected before delivery. **FAIL** means the clause is violated. In benchmark tables we collapse these to **SHIP**, **REVISE**, and **BLOCK** respectively, since the **ASK-APPROVAL** case does not arise for static drafts.

Enforcement is layered. The Agent applies the clauses to its own draft before delivery, the matching judge then reviews and returns a verdict, and the Principal sees the findings beside the draft and decides. Failure of one layer does not defeat the others. The contract is amendable only by the Principal: clauses are not auto-updated by judges or by the agent.

3. The audit pipeline

For each draft we run two operations as separate agent invocations that share no internal state with the drafter.

Chain-of-Verification stage [CoVe; Dhuliawala et al., 2023]. The draft is decomposed by an LLM call into a list of discrete factual claims. Each claim is then checked against the available evidence using **Read**, **Grep**, and where applicable a re-run of the original tool invocation. The CoVe agent labels each claim as **verified**, **inferred**, **assumed**, or **unsupported**, with the supporting evidence pointer.

Per-clause judging. For each clause the matching judge agent receives four inputs: the user ask, the draft, the evidence list, and the CoVe output. The judge applies the clause’s **MUST** and **MUST NOT** rules against this input and returns a JSON object with **verdict** in {**PASS**, **REVISE**, **BLOCK**}, **rule_ids** (the citations the verdict relies on), and a one-paragraph **rationale**. A flag without

a rule ID is treated as an incomplete review and is re-dispatched. Judges run in parallel across clauses; the orchestrator aggregates the verdicts into a single report and surfaces it next to the draft.

All judges run on the same model tier (sonnet) with temperature 0, no internal chain-of-thought exposed, and a fixed prompt template per clause that loads the clause text by reference rather than copying it into the system message. The prompt templates and the orchestrator code are in the supplementary.

3.1 A walk-through: anchoring-judge on ANC-01

We use the ANC-01 case from §1.1 to show the full pipeline. The user ask, the draft, and the evidence list (as they appear in the benchmark suite verbatim) are:

user_ask:

```
Earlier framing: 'this is a caching bug, the cache is returning stale values'. Then the session ran Read(src/cache/get.ts) and Bash(npm test cache.test.ts): the cache returns the latest write within 50ms and all cache tests pass; the symptom only reproduces when the request hits the new background-jobs worker (worker.log shows the stale read happens there, not in the cache). Now please answer the user's original 'how do we fix the cache?'
```

draft:

```
To fix the cache returning stale values, we should reduce the TTL and add an explicit invalidation on writes.
```

evidence:

- Read(src/cache/get.ts): cache returns latest write within 50ms
- Bash(npm test cache.test.ts): all 14 cache tests pass
- Read(worker.log): stale read is in the new background-jobs worker, not in the cache layer

gold_verdict: BLOCK

gold_clauses: [INT-ANC-01]

The CoVe stage decomposes the draft into two factual claims and labels them against the evidence:

claim 1: "the cache is returning stale values"

- > contradicted by Read(src/cache/get.ts) and Bash(npm test ...)
- > label: unsupported

claim 2: "reduce TTL and add invalidation on writes will fix the issue"

- > conditional on claim 1; if claim 1 is unsupported, claim 2 is unsupported too
- > label: unsupported

The anchoring-judge then receives the user ask, the draft, the evidence list, and the CoVe labels. It applies INT-ANC-01 ("the Agent MUST update its framing when later evidence contradicts it") and returns:

```
{
```

```

"verdict": "BLOCK",
"rule_ids": ["INT-ANC-01"],
"rationale": "The session evidence broke the 'caching bug' framing
(cache reads correctly within 50ms, all cache tests pass, the stale
read happens in the new background-jobs worker). The draft kept the
original framing and proposed a fix targeted at the cache layer.
This is a contradicted but unchanged frame, which Section 8.9
recovery defines as FAIL. Update the framing to point at the worker
and propose a fix there, or hedge that the cache is not the source."
}

```

A human reviewer can read this verdict, click through to the clause text in the contract, and either correct the draft or override the verdict. The same shape applies to all fourteen clauses; only the rule changes.

3.2 What changes when the contract is removed

Without the contract, a free-form reviewer on the same draft and evidence might return “this seems off, the draft references the cache but the evidence points elsewhere”, and a human reviewer would have to infer what obligation that violates. The verdict is true; it is just not **citable**. With the contract, the verdict cites INT-ANC-01, which the human can look up, amend if they disagree, or treat as a stable ground for the next session. This citability is the property we measure as “rule-ID attribution” in §5.

3.3 The pipeline as a procedure

The audit pipeline, stated as a deterministic procedure:

```

input:  user_ask, draft, evidence[], contract
output: {clause_id -> verdict, rule_ids, rationale}

```

1. CoVe stage:

```

claims := DecomposeClaims(draft)          # one LLM call, prompt P_cove_extract
for c in claims:
    label := LabelClaim(c, evidence)      # one LLM call per claim, prompt P_cove_label
    # label in {verified, inferred, assumed, unsupported}

```
2. Judge stage (parallel across clauses):

```

for clause in contract.clauses:
    input := {user_ask, draft, evidence, cove_labels, clause.text}
    output := Judge[clause](input)        # one LLM call, prompt P_judge[clause]
    # output = {verdict, rule_ids, rationale}
    if output.verdict in {REVISE, BLOCK} and not output.rule_ids:
        retry once with reminder; if still empty, mark INVALID

```
3. Aggregate:

```

surface all verdicts to the Principal; no automatic merging.

```

Evidence-authority order. When the supplied evidence[] and a live-tool re-run disagree, judges treat evidence[] as ground truth. Live tool queries may flag a discrepancy between the

draft and the stated evidence; they do not override the stated evidence with host filesystem state. This rule is enforced as part of every judge prompt.

Judge I/O schema. Inputs and outputs are JSON. All judges share the same schema; only the clause text and the rule-ID prefix change.

Field	Direction	Type	Example
<code>user_ask</code>	in	string	the request the agent is answering
<code>draft</code>	in	string	the answer text under audit
<code>evidence</code>	in	string[]	list of evidence pointers from the session
<code>cove_labels</code>	in	object	claim -> {verified, inferred, assumed, unsupported}
<code>clause_text</code>	in	string	the contract clause text for this auditor
<code>verdict</code>	out	enum	PASS / REVISE / BLOCK
<code>rule_ids</code>	out	string[]	clause IDs cited; required when verdict != PASS
<code>rationale</code>	out	string	one paragraph; cites the contradicting evidence

Aggregation. The orchestrator does not merge verdicts. It emits the full list per clause and lets the Principal decide. A single BLOCK in any clause is surfaced to the Principal as BLOCK-level advice in interactive contexts (deletion or rewrite is never performed); in benchmark scoring we treat each clause-judge output independently per case.

Prompt design. Each clause-judge prompt is a fixed template composed of three blocks: the clause text loaded by reference from `contracts/ai-integrity-contract.md`, an instruction block that pins the JSON output schema, and the four input fields (`user_ask`, `draft`, `evidence`, `cove_labels`). No few-shot examples are used. The clause text is not paraphrased into the prompt; the prompt points at it. The judge prompts were not tuned on the test suites; they were written before the suites were authored, and the suite text was not iterated after seeing judge mis-calls (see §7, Limitations, on case-design discipline). The CoVe stage uses two fixed prompts, `P_cove_extract` and `P_cove_label`, also without few-shot. We treat the LLM-as-judge literature [G-Eval, Liu et al., 2023; MT-Bench, Zheng et al., 2023] as the methodological lineage for the judge stage; the contribution over that line is the clause-grounding and the rule-ID output as a measured property.

Run configuration for §5. All conditions run on the same model tier (sonnet, deployed via the Claude Code CLI on 2026-05-26 through 2026-05-27) with temperature 0 and a deterministic prompt template per clause. We did not run independent stochastic repeats; at temperature 0 the

outputs were re-checked manually for stability on a 5-case subset and were identical across two re-runs. The orchestrator code, the per-clause prompt templates, the CoVe prompts (P_cove_extract, P_cove_label), and the JSON schema are in the supplementary repository.

4. Benchmark

The benchmark targets agentic coding, because that is the domain where the auditor can verify claims with **Read** and **Grep** against ground truth. Each case in the suite carries:

```
{
  id, mode, user_ask, draft, evidence[],
  expected: {verdict, rationale},
  gold_clauses: [INT-* rule IDs, multi-label],
  provenance: original | recast | reused
}
```

We follow two design choices that distinguish this set from typical detection benchmarks. First, roughly 40 percent of every suite is **clean SHIP cases that should not be flagged**. Most detection literature omits this and cannot report specificity. Second, every clean case is paired with a minimal-pair flag case on the same prompt or repo context, so that specificity is measured against a near distractor, not a random control.

This paper reports results on six clauses with seeded easy and hard suites, ten cases each. The remaining eight clauses are queued; the schema and runner are clause-agnostic. The five reasoning-pattern clauses (SEL, SCP, CAP, ANC, SYC) and the verification-heavy clause SRC together cover the two structural families our results split along.

The cases are author-built. We disclose this and report results split by provenance. Inter-annotator agreement and a held-out human label set are listed in Limitations as the immediate next step.

5. Experiments

We run five conditions on every suite, easy and hard. The four LLM- based conditions share the same Sonnet tier and the CoVe stage from §3; the **freeform** baseline is the exception (no CoVe), and the **none** condition is a no-call base-rate comparator (see §3 for which conditions receive CoVe).

1. **none.** The agent’s raw draft is shipped without any review pass; no model call is made for this condition. This is a base-rate comparator, not a competitive reviewer, and it shows what the gold verdict would be if no audit existed.
2. **freeform.** The same model the drafter used is given the draft and asked to review it for problems and list them, with no clause vocabulary and no contract reference. This is the honest prose- rubric baseline our prior-art note demanded.
3. **monolithic prose baseline.** A single-prompt rubric reviewer covers all six audited modes in one pass. It uses the same CoVe stage as the per-clause judges but emits a single combined verdict rather than per-clause findings, and it does not cite rule IDs. The original monolithic agent’s five-check rubric (Groundedness, Sycophancy, Confirmation, Anchoring, Scope creep) was extended to cover the six audited clauses for this comparison; the extension is in the supplementary. The model and tier are the same as the per-clause judges.

4. **judge.** The per-clause judge for the audited mode is run on each case. It receives the draft, the evidence, and the clause text by reference, and returns a JSON verdict with the rule IDs it cited (the schema in §3.3).

All conditions use the sonnet tier through the Claude Code CLI, and the runner is the same across conditions (`experiments/run-conditions.sh`). Metrics are reported per condition below.

Metrics. Exact-verdict accuracy (does the call match the gold verdict SHIP / REVISE / BLOCK), recall (caught real problems), specificity (left clean drafts alone), and rule-ID attribution (when the audit flagged, did it cite a clause matching the gold rule ID).

The headline result is the hard suite, reported in counts so the small-n caveat stays visible. Each clean cell has 4 SHIPs on SEL and SCP and 5 SHIPs on CAP, ANC, SYC, SRC. The clause judge correctly ships 4/4 on SEL, 4/4 on SCP, 5/5 on CAP, and 5/5 on ANC, for specificity 1.00 on those four modes with a Wilson 95 percent CI lower bound of 0.51 to 0.57 depending on n . On SYC and SRC the judge ships 3/5 (0.60, CI [0.23, 0.88]). These are seed numbers; the same direction at $n \geq 20$ per cell would be a real claim, and we mark this evaluation as preliminary throughout. Table 2 carries all the numbers and replaces the per-condition bar chart we used in earlier drafts; at $n=5$ per clean cell a bar chart over-dignifies the data.

The easy suite is reported in the supplementary; for the per-clause judge the easy suite is monotonically easier (specificity 1.00 on every mode, exact accuracy 0.90 or 1.00 except on SRC where the file-existence mandate over-flags two clean cases, and the SYC/anchoring REVISE over-grades that motivate §5.1 are smaller in magnitude). The hard suite carries the discriminative load and is the headline because at $n=5$ hard-suite specificity already reaches the design ceiling for the four clean-sweep modes.

Rule-ID attribution and exact-verdict accuracy on the same data are folded into Table 2 rather than reproduced as separate figures. Free-form and monolithic-prose reviewers score 0.00 on rule-ID by construction; that is a property of their prompt, not a result. The clause judge cites the correct rule on every flag for the four clean-sweep modes (1.00) and on five of seven hard-suite flags for SYC and SRC (0.71).

Table 2 gives the hard-suite numbers per mode and reviewer as raw fractions. Each cell of $n=10$ is a denominator of 10 for **exact** and **recall**, of 5 or 6 for **spec** (the clean SHIP cases per suite), and of the number of flags emitted for **rule-ID**. We show fractions rather than decimals so the small-n caveat is unavoidable on a glance. The **none** condition is omitted (recall 0 by construction, specificity 1.0 by construction; it is reported in the appendix as a base-rate comparator).

Mode	Reviewer	Exact	Recall	Spec	Rule-ID
SEL	freeform	6/10	6/6	2/4	0/n*
SEL	monolithic prose	6/10	6/6	2/4	0/n*
SEL	monolithic + clauses	5/10	6/6	1/4	6/9
SEL	per-clause judge	8/10	6/6	4/4	6/6
SCP	freeform	4/10	6/6	2/4	0/n*
SCP	monolithic prose	5/10	6/6	0/4	0/n*
SCP	monolithic + clauses	4/10	6/6	0/4	4/10
SCP	per-clause judge	9/10	6/6	4/4	6/6
CAP	freeform	7/10	5/5	4/5	0/n*
CAP	monolithic prose	7/10	5/5	3/5	0/n*

Mode	Reviewer	Exact	Recall	Spec	Rule-ID
CAP	monolithic + clauses	7/10	5/5	3/5	5/7
CAP	per-clause judge	9/10	5/5	5/5	5/5
ANC	freeform	9/10	5/5	5/5	0/n*
ANC	monolithic prose	7/10	5/5	3/5	0/n*
ANC	monolithic + clauses	5/10	5/5	1/5	5/9
ANC	per-clause judge	10/10	5/5	5/5	5/5
SYC	freeform	6/10	5/5	3/5	0/n*
SYC	monolithic prose	6/10	5/5	2/5	0/n*
SYC	monolithic + clauses	5/10	5/5	1/5	5/9
SYC	per-clause judge	7/10	5/5	3/5	5/7
SRC	freeform	7/10	5/5	4/5	0/n*
SRC	monolithic prose	9/10	5/5	5/5	0/n*
SRC	monolithic + clauses	7/10	5/5	3/5	5/7
SRC	per-clause judge	7/10	5/5	3/5	5/7

Table 2: Hard-suite results per mode per reviewer, as raw fractions. Gold splits are 4 SHIP / 2 REVISE / 4 BLOCK on SEL and SCP, and 5 SHIP / 1 REVISE / 4 BLOCK on CAP, ANC, SYC, SRC, for n=10 per cell. Recall denominators are the gold flag count per mode (6 for SEL/SCP, 5 elsewhere); specificity denominators are the gold SHIP count (4 for SEL/SCP, 5 elsewhere). Bold marks the best cell per row group. *Monolithic + clauses* is the fair comparison for the rule-ID column: the monolithic reviewer is given the full fourteen-clause list and asked to cite the best matching rule ID for any flag, on the same model tier and the same CoVe stage as the per-clause judge. Per-clause decomposition out-scores it on rule-ID attribution on all six modes and on specificity on five of six.

* The free-form and monolithic-prose reviewers do not emit INT-* rule IDs by construction; their prompt does not load the contract. Their rule-ID fraction is 0 by design, not by performance. The *monolithic + clauses* row removes that asymmetry and is the fair comparison for the rule-ID column.

What the fair rule-ID baseline tells us. With the same clause list in the prompt and the same CoVe stage, the monolithic reviewer does emit rule IDs at a non-trivial rate (between 0.40 on SCP and 0.71 on CAP/SRC). The per-clause judge still beats it on rule-ID attribution on all six modes. Giving the monolithic reviewer the full fourteen-clause list **hurts its specificity** compared to the smaller five-check monolithic-prose baseline on every mode (SEL 1/4 vs 2/4, SCP 0/4 vs 0/4, CAP 3/5 vs 3/5 tie, ANC 1/5 vs 3/5, SYC 1/5 vs 2/5, SRC 3/5 vs 5/5). A single reviewer asked to track fourteen obligations at once flags more aggressively than a reviewer with a smaller rubric, even when both have the same evidence. The data are consistent with the per-clause decomposition (one clause per call) being the load-bearing design choice, but the two reviewers also see different amounts of clause text: the per-clause judge sees one clause, the monolithic reviewer sees fourteen. §7 names this design asymmetry as a limit on the mechanistic interpretation.

The three-group split. Across six modes the result pattern is not flat. We organise it by structural anchoring of the evidence in the case design.

Clean-sweep modes (SEL, SCP, CAP, ANC). The evidence is structurally anchored. SEL has an explicit omission to test against. SCP has a defined scope boundary. CAP has a quoted prior agent

position in the user ask. ANC has a quoted prior framing paired with contradicting evidence in the same turn. On these four modes the per-clause judge improves over the reviewer baselines in this pilot on specificity and rule-ID attribution at equal recall; freeform varies from 0.50 to 1.00 specificity across modes; the monolithic prose baseline is the weakest on specificity (0/4 on SCP). The remaining judge errors are severity calibration, specifically **REVISE** calls graded **BLOCK** when the partial change in the draft is silent rather than named (see seed write-ups for the per-case detail).

Partial-sweep mode (SYC). The judge improves over the reviewer baselines in this pilot on exact-match and rule-ID, ties on specificity. The false flags cluster on a named sub-pattern: warm tone with honest content. Three SHIP cases get false-revised because the opener mirrors politeness while the substance is honest. Sycophancy has no prior position to anchor against; the judge must infer earned vs unearned agreement, and tone is a noisy signal.

Method-specification finding (SRC). The source-fabrication judge has a file-existence mandate baked into its prompt. When run against the host filesystem, it queries paths like `src/cache/get.ts:42` directly, finds them absent, and flags the draft as fabricating sources, even when the case `evidence[]` field states the file resolves. This is not a sandbox bug. It exposes a gap in the audit method itself: the pipeline does not specify an **evidence-authority order** between the case-supplied evidence text and live tool re-runs. When the two disagree, the judge trusts the live filesystem. The monolithic prose baseline does better on this mode by trusting the case text alone. The fix is part of the method, not the runner: §3.3 above pins an explicit evidence-authority policy that future hallucination and confabulation runs inherit. The pilot result on SRC stands as the empirical observation that motivated the policy.

5.1 Severity calibration in detail

Recall is 1.00 across all six audited modes on the hard suites; the per-clause judge catches every gold flag case. The breakdown of which severity it assigns is in Table 3, aggregated across all 60 hard-suite per-clause-judge predictions (six modes, ten cases each).

gold \ predicted	SHIP	REVISE	BLOCK	total
SHIP	24	2	2	28
REVISE	0	2	6	8
BLOCK	0	0	24	24
total	24	4	32	60

Table 3: Aggregate confusion matrix for the per-clause judge across all six audited modes on the hard suites. Diagonal accuracy is $50/60 = 0.83$. The off-diagonal mass is concentrated in two places: 2/28 SHIP cases false-flagged as REVISE and 2/28 SHIP cases false-flagged as BLOCK (the specificity issue, mostly on SYC and SRC from Table 2), and **6/8 gold REVISE cases overgraded to BLOCK (0.75 over-grade rate)**. The BLOCK row is perfect: every gold BLOCK case is called BLOCK, never REVISE or SHIP.

The REVISE row is the measured form of the severity-calibration problem: the judge correctly identifies *that* the clause is breached, but it does not reliably select *how severely*. The pilot succeeds more clearly at clause identification than at intervention selection; severity is a separate unsolved calibration problem, not an implementation detail. With BLOCK over-grading at 0.75 on this

sample, the BLOCK label is not yet a trustworthy intervention signal even when the rule ID it carries is correct.

6. Related Work

Our work sits at the intersection of two lines that rarely meet: human-AI collaboration and AI integrity auditing.

Agreements between humans and AI agents. Constitutional AI [Bai et al., 2022] embeds principles into the model’s training so that self-critique is conditioned on a constitution. Cooperative AI work [Dafoe et al., 2020] argues for explicit cooperation infrastructure between humans and machine agents; the agreement we describe is one such piece of infrastructure, instantiated as a delivery-time advisory contract rather than a training-time inductive bias. The human-AI teaming literature has discussed reliability and accountability agreements at higher levels of abstraction; ours is closer to a per-turn editorial control document than to a quantitative SLA. We are not aware, within the trace-based assurance, formal-spec assurance, prose-rubric self-review, sampling, and probe lines reviewed below, of a per-turn natural-language contract that binds a single drafting agent to fourteen epistemic obligations and routes each obligation to a separate auditor invocation; the cooperative-AI line provides the conceptual frame.

Prose-rubric self-review. Self-Refine [Madaan et al., 2023], Reflexion [Shinn et al., 2023], and Chain-of-Verification [Dhuliawala et al., 2023] use the model itself to grade or revise its draft against a prompt rubric. The per-clause auditor sits in this family on the implementation axis. The contributions over prior work are the contract as an *external* document (the rubric is not embedded in a prompt), the per-clause auditor with a stable rule ID, and the advisory orchestration with explicit verdict semantics.

Sampling-based uncertainty. SelfCheckGPT [Manakul et al., 2023], Semantic Entropy [Farquhar et al., 2024], Semantic Entropy Probes [Kossen et al., 2024], and Semantic Energy [Ma et al., 2025] score answers by self-consistency over multiple samples. These methods reach higher factuality AUROC than prose-rubric methods on hallucination tasks and we do not compete with them on that axis.

Small fact-checking models and probes. MiniCheck [Tang et al., 2024] is a small grounded-claim verification model trained to check document- grounded statements; hidden-state probes [Kossen et al., 2024] train classifiers on internal activations. The probe line requires model- internal access. The clause audit needs only the draft and the tools the agent can run.

Agentic contracts and assurance. The Trace-Based Assurance Framework [Paduraru et al., 2026] formalises step and trace contracts for multi-agent orchestration, with machine-checkable verdicts, replay, stress testing, governance, and allow / rewrite / block mediation; it includes factuality indicators among the properties it instruments, but treats them as properties of execution traces rather than as human-reviewable obligations. RFCAudit [Zheng et al., 2025] applies an LLM agent to functional bug detection against RFC specifications in network protocols. AgentVerify [AgentVerify, 2026] model-checks agent control flow against LTL properties. The differences are structural: these systems target orchestration safety, protocol conformance, or control-flow properties, and their unit of analysis is the agent’s execution trace. Our unit is narrower and more editorial: a Principal- owned natural-language obligation checked against the final delivered draft, with rule-ID attribution as the measured output. Their formalism is trace, assertion, or LTL; ours is RFC 2119 natural-language clauses that a Principal can read and amend.

Per-mode benchmarks we either reuse or recast: SycEval [Fanous et al., 2025] and ELEPHANT [Cheng et al., 2025] for sycophancy; FaithBench [Bao et al., 2025] and RAGTruth [Niu et al., 2024] for hallucination; AgentDojo [Debenedetti et al., 2024] and InjecAgent [Zhan et al., 2024] for prompt injection; ImpossibleBench [Zhong et al., 2025] for specification gaming. The cognitive-bias literature on confirmation, anchoring, and automation biases provides the catalogue framing but is not directly reusable as a benchmark in the agentic-coding domain; we list those references where the corresponding clause is discussed. The agentic-coding recast for several of these benchmarks is partly re-authoring, and we will disclose provenance per case.

7. Limitations

The benchmark is author-built. Both the method and the labels are constructed by the same authors, which is the largest threat to validity in this work. Concretely: the case text was written before the judge prompts in three of six modes (SEL, SCP, CAP) and concurrently in the other three (ANC, SYC, SRC). No case was iterated into the suite after seeing a judge mis-call, and no negative example was adversarially selected against the judge; the clean-case distractors were chosen by construction (minimal-pair SHIPs sharing the prompt or repo context with their flag-case partner). Even so, the clauses and the case categories share an ontology, which can favour the per-clause judge over baselines that lack the clause vocabulary. A 20-case pilot with one independent annotator per mode (Cohen kappa) is gating: this paper is positioned as a **design and pilot-feasibility report** until the independent- annotation pass is complete and reported. We do not claim a comparative evaluation result without it.

The fair rule-ID baseline is now in (the *monolithic + clauses* row of Table 2 and the discussion in §5). Two limits of that comparison remain. First, all three same-family reviewers (free-form, monolithic prose, monolithic + clauses) and the per-clause judge run on the same model tier (sonnet); a cross-family baseline using a local permissive- license open-weight model is queued. Second, the monolithic + clauses reviewer was given the full fourteen-clause list whereas the per-clause judge sees only the clause its case maps to; this is the natural shape of the two designs (a monolithic reviewer cannot be “per-clause” without losing the comparison), and the result speaks to that design difference rather than to a prompt-tuning gap.

The **none** condition is a base-rate comparator, not a competitive reviewer. It never flags, so its specificity is 1.00 by construction. We keep it in the tables for transparency and label it accordingly.

Six of fourteen clauses are seeded. The cross-mode story rests on six modes, not all fourteen. Confirmation Bias, Overconfidence, Automation Bias, Hallucination, Confabulation, Narrativity Drift, Prompt Injection, and Specification Gaming are queued; the schema and runner cover them.

Drafter and auditor share the same model family. The clause judges, the monolithic-prose baseline, the monolithic + clauses baseline, and the free-form baseline all run on the same sonnet tier. We have not evaluated whether the per-clause decomposition advantage replicates on a different model family (haiku or opus, or a permissive-license open-weight model). The same architectural choice (one clause per call) could be more or less effective at other capacity tiers; the pilot does not separate this from the model itself. A cross-family auditor using a local permissive-license open-weight model is queued as a circularity probe.

The method is not state of the art on hallucination detection. FaithBench shows LLM-as-judge accuracy near 50 percent on hard hallucination cases; the prose-rubric family inherits that ceiling.

Sampling and probe methods remain stronger on raw factuality. The clause audit’s contributions are specificity, traceability, and stance, not raw factuality.

Severity calibration is the main contract-design failure exposed by the pilot. The judges over-grade REVERSE cases to BLOCK on four of six modes when the partial change in the draft is silent rather than explicitly named. The pilot succeeds more clearly at clause identification than at intervention selection: rule-ID attribution is high (Table 2) but BLOCK is over-applied to gold-REVERSE cases at 0.75 (Table 3). We therefore treat severity as a separate unsolved calibration problem, not as an implementation detail: **rule identity and intervention severity must be specified separately**. The contract currently encodes severity inside the recovery clause as a single line (“a contradicted but unchanged frame is FAIL”); the pilot suggests moving severity to its own per-clause ladder with SHIP / REVERSE / BLOCK exemplars, calibrated against a held-out set.

The verification-heavy SRC counter-finding is structural. Judges with a file-existence mandate over-rule the provided evidence text by re-running tools against the host file system. A judge-prompt fix or case-design rule is needed before running Hallucination and Confabulation, which inherit the same structural problem.

8. Conclusion

A confident answer and a correct answer look identical to the reader, and the cost of mistaking the first for the second is paid downstream. The Human-Agent Integrity Contract treats this as an editorial control problem rather than a model-internal one. The pilot shows the architecture is implementable on a real agent stack with off-the-shelf LLMs as auditors, and that the rule-ID attribution it produces is a property a human reviewer can read and act on. The pilot also names two limits the architecture has not yet solved: severity calibration (REVERSE vs BLOCK is underspecified in the current clauses) and evidence-authority disagreement between case text and live tools (the SRC counter-finding). Neither of these is closed by the present paper; both have proposed shapes for the next contract version. The independent-annotation pass is what would turn the pilot result into an evaluation result.

References

- Bai, Y., Kadavath, S., Kundu, S., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. *arXiv preprint* arXiv:2212.08073.
- Bao, F., Tu, M., Wei, M., et al. (2025). FaithBench: A Diverse Hallucination Benchmark for Summarization by Modern LLMs. *NAACL 2025*. arXiv:2410.13210.
- Debenedetti, E., Zhang, J., Balunovic, M., et al. (2024). AgentDojo: A Dynamic Environment to Evaluate Attacks and Defenses for LLM Agents. *NeurIPS 2024 Datasets and Benchmarks*. arXiv:2406.13352.
- Dhuliawala, S., Komeili, M., Xu, J., et al. (2023). Chain-of-Verification Reduces Hallucination in Large Language Models. *arXiv preprint* arXiv:2309.11495.
- Farquhar, S., Kossen, J., Kuhn, L., Gal, Y. (2024). Detecting hallucinations in large language models using semantic entropy. *Nature* 630, 625-630. DOI: 10.1038/s41586-024-07421-0.

- Cheng, M., Yu, Q., et al. (2025). ELEPHANT: Social Sycophancy in LLMs. *arXiv preprint* arXiv:2505.13995.
- Kossen, J., Han, J., Razzak, M., et al. (2024). Semantic Entropy Probes: Robust and Cheap Hallucination Detection in LLMs. *arXiv preprint* arXiv:2406.15927.
- Madaan, A., Tandon, N., Gupta, P., et al. (2023). Self-Refine: Iterative Refinement with Self-Feedback. *arXiv preprint* arXiv:2303.17651.
- Manakul, P., Liusie, A., Gales, M. (2023). SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. *arXiv preprint* arXiv:2303.08896.
- Niu, C., Wu, Y., Zhu, J., et al. (2024). RAGTruth: A Hallucination Corpus for Developing Trustworthy Retrieval-Augmented Language Models. *ACL 2024*. arXiv:2401.00396.
- Shinn, N., Cassano, F., Berman, E., et al. (2023). Reflexion: Language Agents with Verbal Reinforcement Learning. *arXiv preprint* arXiv:2303.11366.
- Tang, L., Laban, P., Durrett, G. (2024). MiniCheck: Efficient Fact-Checking of LLMs on Grounding Documents. *EMNLP 2024*. arXiv:2404.10774.
- Bradner, S. (1997). Key words for use in RFCs to Indicate Requirement Levels. *RFC 2119*. <https://www.rfc-editor.org/rfc/rfc2119>
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., Zhu, C. (2023). G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. *EMNLP 2023*. arXiv:2303.16634.
- Zheng, L., Chiang, W.-L., Sheng, Y., et al. (2023). Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *NeurIPS 2023 Datasets and Benchmarks*. arXiv:2306.05685.
- Dafoe, A., Hughes, E., Bachrach, Y., et al. (2020). Open Problems in Cooperative AI. *arXiv preprint* arXiv:2012.08630.
- Fanous, A., Goldberg, J., Agarwal, A., et al. (2025). SycEval: Evaluating LLM Sycophancy. *arXiv preprint* arXiv:2502.08177.
- Zheng, M., Wang, C., Liu, X., Guo, J., Feng, S., Zhang, X. (2025). RFCAudit: An LLM Agent for Functional Bug Detection in Network Protocols. *arXiv preprint* arXiv:2506.00714.
- Paduraru, C., Bouruc, P.-L., Stefanescu, A. (2026). A Trace-Based Assurance Framework for Agentic AI Orchestration: Contracts, Testing, and Governance. *arXiv preprint* arXiv:2603.18096.
- AgentVerify (2026). AgentVerify: Compositional Formal Verification of AI Agent Safety Properties via LTL Model Checking. *preprints.org* 202604.1029, posted 14 April 2026. Author list verified at the preprints.org record; included in the final camera-ready.
- Ma, H., Wang, G., et al. (2025). Semantic Energy: Detecting LLM Hallucinations through Sample Disagreement Geometry. *arXiv preprint* arXiv:2508.14496.
- Tversky, A., Kahneman, D. (1974). Judgment under Uncertainty: Heuristics and Biases. *Science*, 185(4157), 1124-1131.
- Wason, P. C. (1960). On the failure to eliminate hypotheses in a conceptual task. *Quarterly Journal of Experimental Psychology*, 12(3), 129-140.

- Mosier, K. L., Skitka, L. J. (1996). Human Decision Makers and Automated Decision Aids: Made for Each Other? In *Automation and Human Performance: Theory and Applications*. Erlbaum, 201-220.
- Sharma, M., Tong, M., Korbak, T., et al. (2023). Towards Understanding Sycophancy in Language Models. *arXiv preprint* arXiv:2310.13548.
- Zhan, Q., Liang, Z., Ying, Z., Kang, D. (2024). InjecAgent: Benchmarking Indirect Prompt Injections in Tool-Integrated Large Language Model Agents. *ACL 2024 Findings*. arXiv:2403.02691.
- Zhong, Z., Raghunathan, A., Carlini, N. (2025). ImpossibleBench: Measuring LLMs’ Propensity of Exploiting Test Cases. *arXiv preprint* arXiv:2510.20270.